

Program Design & Data Structures (Course 1DL201)

Uppsala University Autumn 2014/Spring 2015

Homework Assignment 1: Computational Linguistics

Prepared by Dave Clarke

Lab: Friday, 21 November, 2014
Submission Deadline: 18:00, Friday, 28 November, 2014
Lesson: Thursday/Friday, 4/5 December, 2014
Resubmission Deadline: 18:00, Friday, 12 December, 2014

Computational Linguistics

Computational Linguistics is the discipline that studies natural language from a computational perspective by using computational models to analyse vast volumes of text in order to understand how language is used, among other things. Corpus Linguistics is one particular branch of Computational Linguistics that heavily uses computational techniques to understand the meaning of words from the contexts in which they are used. Specifically, Corpus Linguistics is interested in *collocations*, sequences of words that appear together more often than one would expect by chance. For example, one would expect to see *strong tea* in texts more often than *powerful tea*, yet *powerful computer* would appear more often than *strong computer*—ultimately stating something interesting to Computational Linguists about the difference in meaning of *strong* and *powerful*.

Goal

The goal of this assignment is to write a few Haskell functions for a computational linguistics toolbox.

The following are the core data types to be used in your program:¹

```
type Document = [String]
type WordTally = [(String, Int)]
type Pairs = [(String, String)]
type PairsTally = [((String, String), Int)]
```

Type `Document` represents documents as lists of words—one can assume that raw documents have been preprocessed and the words have been extracted. The `WordTally` data type can be used to record word counts. Data type `Pairs` is used for lists of words that are *adjacent* to each other in

¹The keyword `type` introduces a synonym for an existing type. For instance, declaration `type Document = [String]` allows the programmer to use `Document` interchangeably with `[String]`.

the text. Finally, `PairsTally` is similar to `WordTally` except that it is used to record the number of occurrences of pairs of words.

The functions to be implemented (described in the next section) take a document and compute word counts, pairs of words that occur together, and simple statistics about those. The first couple of chapters of Jane Austin's *Pride and Prejudice* are available (variable `austin` in module `PandP`) to experiment with.

Work to be Done

Download the files from the student portal: `PandP.hs` contains preprocessed text from *Pride and Prejudice*, and `CompLing.hs` is for your solution. `CompLing.hs` also includes a number of tests to help develop your implementation.

You are required to implement the following functions (empty implementations are provided in the file `CompLing.hs`):

1. Function `wordCount :: Document -> WordTally` computes a tally of all the distinct words appearing in the document. For example, the text "A rose is a rose is a rose." is encoded in Haskell as the list of words `["a", "rose", "is", "a", "rose", "is", "a", "rose"]`. The tally for this document is:

- a — 3
- rose — 3
- is — 2

This result could be represented in Haskell as follows, though the order of elements in the list is not specified:

```
[("rose", 3), ("a", 3), ("is", 2)]
```

2. Function `adjacentPairs :: Document -> Pairs` results in a list of all adjacent pairs of words appearing in the document.

For instance

```
adjacentPairs ["time", "for", "a", "break"]
= [("time","for"),("for","a"),("a","break")]
```

3. Function `pairsCount :: Pairs -> PairsTally` computes a tally of all pairs, such as those computed by `adjacentPairs`.

For instance, `pairsCount [("big","bear"),("bear","big"),("big","dog")]` would result in the tally:

- big, bear — 2
- big, dog — 1

Note that we do not care about the order of words in a pair. For instance, we consider `("big","bear")` and `("bear","big")` to both represent the same pair of words, so they must not be counted separately.

How this tally is represented in Haskell is a design decision, meaning that either:

```
[(("bear","big"),2), (("big","dog"),1)]
```

or

```
[(("big","bear"),2), (("big","dog"),1)]
```

or any reordering of these, are all valid ways of representing the tally.

4. Function `neighbours :: PairsTally -> String -> WordTally` takes a tally of pairs, such as computed by the `pairsCount` function, and a word and gives all the words that appear with that word in the tally of pairs along with the number of occurrences.
5. Function `mostCommonNeighbour :: PairsTally -> String -> Maybe String` returns the word that occurs most frequently with a given word, based on a tally of pairs. The `Maybe` data type is used to represent the result:
 - If the word does not appear in the tally, the result should be `Nothing`
 - If the word appears in the tally, and the neighbour with the largest count is `"foo"`, then the result is `Just "foo"`. If more than one word appears in neighbours with equal highest count, then the result is `Just x`, where `x` is one of those words.

The `Maybe` data type is described in the next section.

The Maybe data type

The `Maybe` data type is used for the result of functions that do not always return a valid result. It is declared in the standard library as

```
data Maybe a = Just a | Nothing
```

A valid result is indicated as `Just x`, for some value `x`, and no valid result is indicated as `Nothing`.

Grading

Grading Your solution is graded on a U/K/4/5 scale based on two components: (1) functional correctness and (2) style and comments.

1. Functional correctness:

Your program will be run on an unspecified number of grading test cases that satisfy all preconditions but also check boundary conditions. Each test case is based on queries to some (possibly empty) document using the functions provided.

We reserve the right to run these tests automatically, so be careful to match exactly the imposed file names, function names, and argument orders—that is, don’t change what we’ve provided for you.

Advice: Run your code in a freshly started Haskell session before you submit, so that declarations that you may have made manually do not interfere when you test the code.

The grade for this component is computed as follows:

- If your solution was submitted by the deadline, your file `CompLing.hs` loads in `ghci` and it passes 70% of the test cases in provided in `CompLing.hs`, you get (at least) a K for functional correctness. Otherwise (including when no solution was submitted by the deadline), you get a U grade for the homework assignment.
- If your program additionally passes at least 80% of the grading test cases, you get (at least) a 4 for functional correctness.
- If your program passes all grading test cases, you get a 5 for functional correctness.

2. Style and comments:

Your program is graded for style and comments according to our *Coding Convention*. The following criteria will be used:

- suitable breakdown of your solution into auxiliary/helper functions,
- function specifications and statements of purpose,
- datatype representation conventions and invariants,
- code readability and indentation,
- sensible naming conventions followed.

The grade for this component is computed as follows:

- If your programs style and comments are deemed a serious attempt at following these criteria, you get (at least) a K for style and comments. Otherwise, you get a U grade for the homework assignment.
- If you have largely followed these criteria, with very few major omissions or errors, you get a 4 for style and comments.
- If you have followed these criteria with at most minor omissions or oversights, you get a 5 for style and comments.

Final Grade

Component grades are converted to a final grade on the usual scale U/3/4/5 as follows:

1. You need to pass both three components (functional correctness and style and comments) in order to pass this assignment.
2. A K grade in either component means that you are required to attend the **lesson** discussing the assignment and to subsequently resubmit the assignment.

After resubmission, your entire assignment will be re-graded. Component grades of K will improve to 3 if you provide a mostly correct solution (cf. the criteria for grade 4); you cannot get a better grade than 3 in an component where you originally got a K. Component grades of 4 or 5 remain unchanged if your resubmission still meets the relevant grading criteria, but may be lowered otherwise (e.g., if your revised program no longer follows the coding convention).

3. Your final grade is the arithmetic mean of the two component grades.

Modalities

- The assignment will be conducted in groups of 2. Groups will be assigned via the student portal.
- Assignments must be submitted via the student portal. Only one solution per group needs to be submitted. Ensure that **both** team members' names appear on all submitted artefacts.

We assume that by submitting a solution you are certifying that it is solely the work of your group, except where explicitly attributed otherwise. We reserve the right to use plagiarism detection tools and point out that they are extremely powerful. You have already been warned about the consequences of cheating and plagiarism.

Good luck!